

## 1) Illustrate snoopy protocols with its approaches.

### Snoopy Protocols and Their Approaches

In multiprocessor systems, maintaining cache coherence is crucial to ensure that multiple caches hold consistent data. One common approach to maintaining cache coherence is using **Snoopy Protocols**. These protocols rely on the ability of each cache to "snoop" on the communication between the other caches and the main memory. Here's a detailed illustration of Snoopy Protocols and their approaches:

#### Basic Concept

- **Snoopy Protocols:** These protocols ensure that all caches that have a copy of a memory block maintain a coherent view of the memory. Each cache controller monitors (or "snoops" on) the bus to determine if it has a copy of a block that is requested or modified by another processor.

#### Types of Snoopy Protocols

1. **Write-Invalidate Protocol**
2. **Write-Update Protocol**

#### Write-Invalidate Protocol

- **Mechanism:** When a processor writes to a block, it invalidates all other copies of that block in other caches.
- **Process:**
  - When a processor wants to write to a cache block, it sends an invalidation signal to other caches.
  - Other caches snoop this signal and invalidate their copies of the block.
  - The writing processor can then update its local copy.
- **Example:**
  - Assume Processor P1 writes to a memory location X.
  - P1 sends an invalidation signal for X.
  - If Processor P2 has X in its cache, it invalidates X.
  - P1 updates X in its cache.
  - Any subsequent read or write by other processors to X will cause them to fetch the updated value from P1 or main memory.
- **Advantage:** Ensures that only one copy of the modified data exists, simplifying consistency.

- **Disadvantage:** Can lead to high invalidation traffic, especially if data is frequently shared and written by multiple processors.

## Write-Update Protocol

- **Mechanism:** When a processor writes to a block, it updates all other copies of that block in other caches.
- **Process:**
  - When a processor wants to write to a cache block, it sends an update signal along with the new data to other caches.
  - Other caches snoop this signal and update their copies of the block with the new data.
- **Example:**
  - Assume Processor P1 writes to a memory location X.
  - P1 sends an update signal for X along with the new value.
  - If Processor P2 has X in its cache, it updates X with the new value.
  - All caches now have the updated value of X.
- **Advantage:** Reduces invalidation traffic and can be more efficient if data is frequently read after being written.
- **Disadvantage:** Can lead to high update traffic, which may be inefficient if updates are frequent and the data is not often read by other processors.

## 2) With neat diagram explain the bus system at board level ,backplane and I/O level

### 1. Board Level Bus System

At the board level, the bus system is integral to the operation of the motherboard, connecting the CPU, memory, and various peripherals.

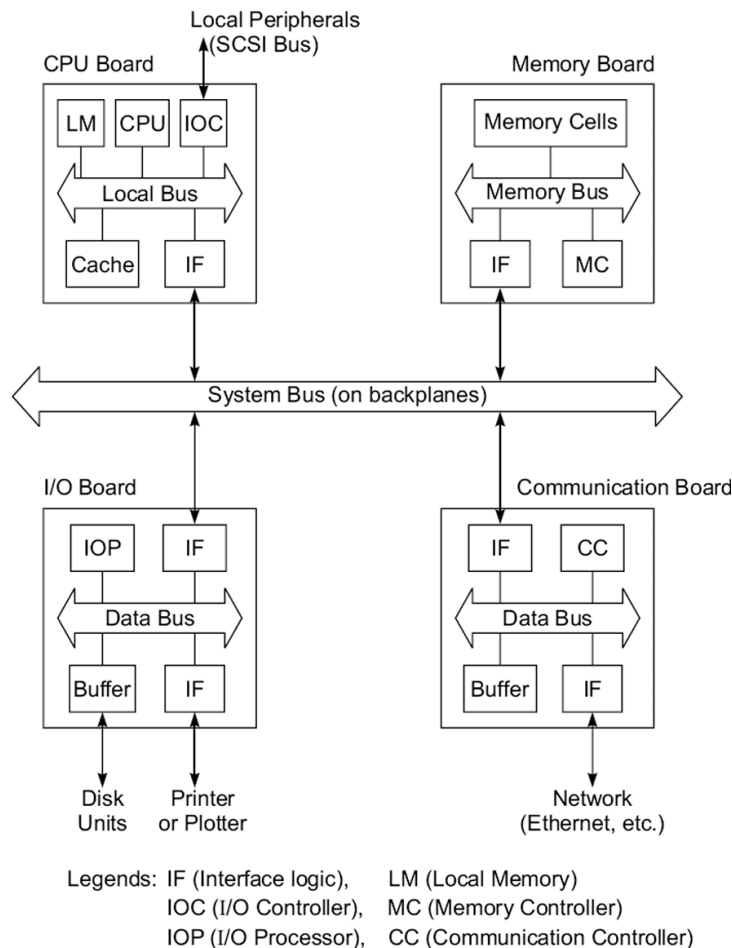
- **Address Bus:** Carries the addresses of data (but not the data itself) between the CPU and memory.
- **Data Bus:** Transports actual data between the CPU, memory, and peripherals.
- **Control Bus:** Carries control signals from the CPU to other components to coordinate various operations.

### Components Connected by the Bus System:

- **CPU (Central Processing Unit):** The main processor that performs all computations and processing tasks.

- **Memory (RAM, ROM, Cache):** Temporary storage for data and instructions that the CPU needs to execute tasks.
- **Peripherals (e.g., USB ports, SATA interfaces):** External devices connected to the motherboard for additional functionality.

The bus system at this level ensures that the CPU can fetch instructions and data from memory and send/receive data to and from peripherals efficiently.



## 2. Backplane Bus System

The backplane bus system is used in systems that require multiple circuit boards to be connected. This is common in servers and high-end workstations where modularity and expandability are critical.

- **Backplane:** A large board that contains multiple slots for connecting various boards (like CPU boards, memory boards, I/O boards). It acts as the backbone of the system.
- **Connections:** The backplane provides the physical and electrical connections between the boards, allowing them to communicate and share power.

**Functionality:**

- **Modularity:** Allows different boards to be easily added or replaced, facilitating upgrades and maintenance.
- **Scalability:** Supports the addition of multiple CPUs, memory modules, and I/O interfaces, making it suitable for high-performance applications.

The backplane bus system is essential for creating flexible and expandable computer systems, particularly in environments that require robust and scalable computing power.

### 3. I/O Level Bus System

At the I/O level, the bus system manages communication between the central processing unit (CPU) and various input/output devices.

- **I/O Bus:** The pathways through which data is transferred between the CPU and peripheral devices. Common I/O buses include PCI (Peripheral Component Interconnect) and USB (Universal Serial Bus).
- **I/O Controllers:** Dedicated hardware that manages data transfer between the CPU and peripheral devices. Examples include disk controllers, network interface controllers, and graphics controllers.

#### Peripheral Devices:

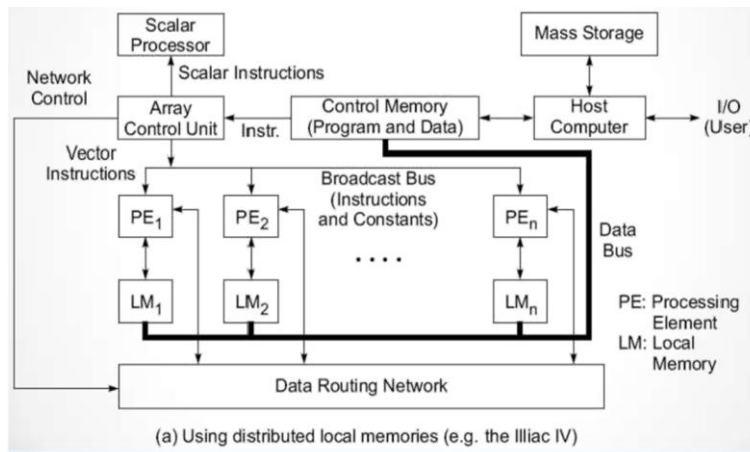
- **Storage Devices (e.g., HDDs, SSDs):** Store data persistently and communicate with the CPU via I/O buses.
- **Input Devices (e.g., Keyboard, Mouse):** Send user inputs to the CPU.
- **Output Devices (e.g., Monitors, Printers):** Receive data from the CPU to provide output to the user.
- **Network Interfaces (e.g., Ethernet, Wi-Fi Adapters):** Enable communication over a network.

### 3) With a neat diagram, describe the implementation models of SIMD.

SIMD (Single Instruction, Multiple Data) models can be distinguished based on their memory distribution and addressing schemes. Two primary models for constructing SIMD supercomputers are the Distributed Memory Model and the Shared Memory Model. Here is a detailed explanation of these models:

#### 1. Distributed Memory Model

In the Distributed Memory Model, each Processing Element (PE) has its own local memory. This model is known for its spatial parallelism, where each PE operates on its own data independently.



### Key Features:

- **Spatial Parallelism:** Exploits parallelism by distributing tasks across multiple PEs, each with its own memory.
- **Array of PEs:** The system consists of an array of PEs, each supplied with local memory, all controlled by an array control unit.
- **Control Unit:** The control unit manages the overall operation. It loads programs and data into the control memory through the host computer and distributes them to the PEs' local memories.
- **Instruction Decoding:**
  - **Scalar Operations:** If the instruction is scalar or involves program control, it is executed by a scalar processor attached to the control unit.
  - **Vector Operations:** If the instruction is a vector operation, it is broadcast to all PEs for parallel execution.
- **Data Distribution:** Partitioned datasets are distributed to the PEs' local memories via a vector data bus.
- **Interconnection Network:** PEs are interconnected by a data routing network, which performs inter-PE data communications like shifting, permutation, and other routing operations.

### Advantages:

- Scalable, as each PE operates independently with its local memory.
- Reduces memory access contention since each PE has its dedicated memory.

### Disadvantages:

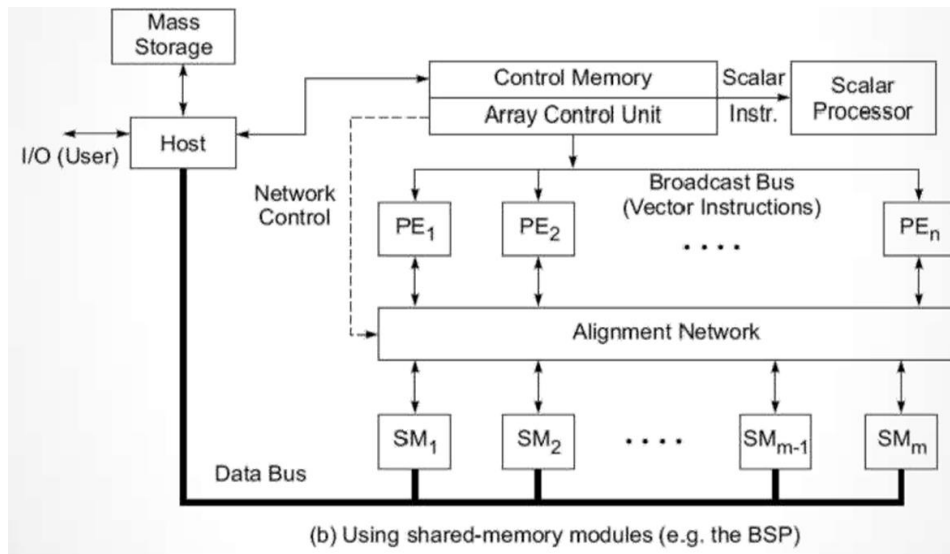
- Communication overhead can be high due to the need for explicit data transfer between PEs.
- Complexity in programming and managing data distribution.

**Example Systems:**

- **MasPar MP-1 and MP-2:** Massively parallel SIMD supercomputers with distributed memory architecture.

**2. Shared Memory Model**

In the Shared Memory Model, all PEs share a common memory space, which allows for more straightforward data access but requires careful management to avoid conflicts.

**Key Features:**

- **Shared Memory:** All PEs access a common shared memory space, controlled by a single control unit.
- **Alignment Network:** An inter-PE memory communication network, called the alignment network, is used to manage memory access. This network ensures that data is correctly aligned and prevents access conflicts.
- **Control Unit:** The control unit coordinates the alignment network and manages the distribution of instructions and data to the PEs.
- **Data Access:** Proper setting of the alignment network is crucial to avoid access conflicts and ensure efficient data transfer among PEs.

**Advantages:**

- Simplifies the programming model, as all PEs can directly access shared data.
- Eliminates the need for explicit data distribution among PEs.

**Disadvantages:**

- Memory access conflicts can occur, requiring sophisticated management techniques.

#### 4) Discuss different vector access memory schemes.

##### Or Explain C-access and S-access organization for m way interleaved memory

Vector processing involves accessing multiple data elements simultaneously, and different memory schemes are used to optimize this access. Here are three primary vector access memory organization schemes: C-Access, S-Access, and C/S-Access.

#### 1. C-Access Memory Organization

The C-Access memory organization involves interleaving memory to allow concurrent and overlapping access to multiple words. This scheme is designed to maximize memory throughput by staggering access cycles.

##### Key Features:

- **Low-Order Memory Structure:** Allows  $m$  words to be accessed concurrently and overlapped.
- **Memory Modules:** The low-order bits of the address select the memory modules, and the high-order bits select the word within each module.
- **Addressing:**  $m = 2^a$ ,  $m = 2^a$ , and  $a + b = n$  (total address length).

##### Operation:

- **Stride of 1:** Successive addresses are latched in the address buffer at the rate of one per cycle. It takes  $m$  minor cycles to fetch  $m$  words, equating to one major memory cycle.
- **Stride of 2:** Successive accesses must be separated by two minor cycles to avoid conflicts, reducing memory throughput by half.
- **Stride of 3:** No module conflict, yielding maximum throughput.
- **General Case:** Maximum throughput ( $m$  words per cycle) is achieved if the stride is relatively prime to  $m$ .

##### Advantages:

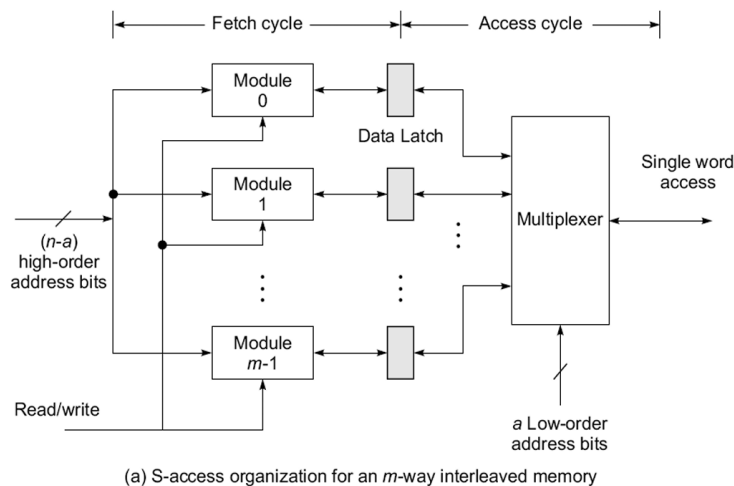
- **High Throughput:** Optimized for strides that are relatively prime to the number of memory modules.
- **Concurrent Access:** Enables concurrent access to multiple words.

##### Disadvantages:

- **Throughput Variability:** Throughput decreases significantly for non-optimal strides.

## 2. S-Access Memory Organization

The S-Access memory organization synchronizes the access to all memory modules, allowing simultaneous access to words across modules.



### Key Features:

- **Synchronized Access:** All memory modules are accessed simultaneously.
- **High-Order Address Bits:** The high-order bits select the same offset word from each module.

### Operation:

- **Memory Cycle:** At the end of each memory cycle,  $m=2^a m = 2^a m = 2^a$  consecutive words are latched.
- **Stride Greater Than 1:** Throughput decreases roughly in proportion to the stride.

### Advantages:

- **Simplicity:** Easy to implement and manage.
- **Concurrent Access:** All modules are accessed at the same time, allowing simultaneous data fetching.

### Disadvantages:

- **Throughput Reduction:** Throughput decreases with increased stride, which can limit performance.

## 3. C/S-Access Memory Organization



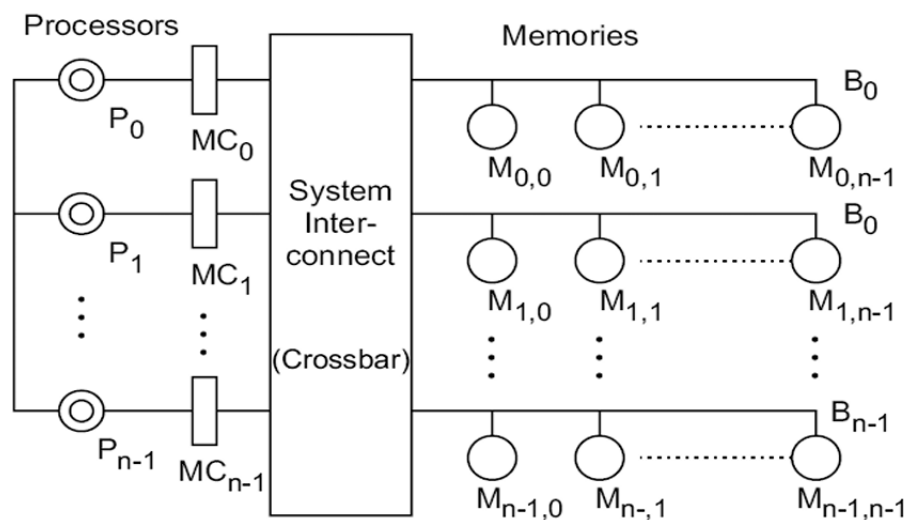
The C/S-Access memory organization combines the features of C-Access and S-Access schemes, providing high bandwidth and parallel pipelined access.

### Key Features:

- **Combination of Schemes:** Utilizes both C-Access and S-Access methods.
- **Access Buses:**  $nmn$  access buses are used with  $mmm$  interleaved memory modules attached to each bus.
- **Interleaving:** The  $mmm$  modules on each bus are interleaved to allow C-Access.

### Operation:

- **Pipelined Memory Accesses:** In each memory cycle, up to  $m \times nm$  \times  $nm \times n$  words can be fetched if the buses are fully used.
- **Parallel Access:** Suitable for vector multiprocessor configurations, providing parallel pipelined access.



### Advantages:

- **High Bandwidth:** Provides high bandwidth access to vector data sets.
- **Parallel Pipelined Access:** Suitable for high-performance vector processing.

### Disadvantages:

- **Complexity:** More complex to implement due to the combination of access methods and multiple buses.

## 5) Illustrate the processor consistency models.

The processor consistency models are covered under the topic of relaxed memory consistency models. These models define how memory operations (loads and stores) performed by one processor appear to other processors. The consistency models vary in their strictness of how operations are ordered and made visible across processors. Here are the primary models discussed in the provided material:

### 1. Sequential Consistency (SC):

- This is the strictest model where the results of execution are as if all operations were executed in some sequential order, and the operations of each individual processor appear in this sequence in the order specified by the program.

### 2. Relaxed Consistency Models:

- These models relax one or more requirements of sequential consistency to allow more flexibility in execution and better performance. They do not provide memory consistency at the hardware level; instead, the programmers must ensure consistency using synchronization techniques. The main types of relaxed consistency models include:
  - **Processor Consistency (PC):**
    - Ensures that writes issued by a single processor are observed by all other processors in the order they were issued, but does not guarantee the ordering of writes from different processors.
  - **Weak Consistency (WC):**
    - Allows both reads and writes to be reordered, but requires that all memory operations must be completed before any synchronization operations (e.g., locks) can be performed.
  - **Release Consistency (RC):**
    - Further refines weak consistency by distinguishing between acquire and release operations. An acquire operation (e.g., acquiring a lock) ensures that all previous writes are visible before it is performed, while a release operation (e.g., releasing a lock) ensures that all previous writes are visible to other processors after it is performed.
  - **Eventual Consistency:**
    - This model ensures that if no new updates are made to a given data item, eventually all accesses to that item will return the last updated value. It is often used in distributed systems and web services where high availability is preferred over immediate consistency.

## 6) Describe Coherence problems in data sharing and process migration.

### Coherence Problems in Data Sharing

#### Data Sharing:

- **Writable Data Sharing:** When multiple processors share writable data, coherence problems can occur. Each processor might cache a copy of the shared data, leading to inconsistencies if one processor updates its copy without informing the others.
- **Cache Coherence:** In a multiprocessor system, each processor may have its private cache. If one processor updates a shared data item in its cache, other processors may still have the old value in their caches. This leads to a situation where different caches have different values for the same data, causing incoherence.
- **Write Through vs. Write Back Policies:**
  - **Write Through:** When a processor writes data to its cache, it also writes the same data to the main memory. This approach can still cause inconsistencies if the other processors' caches are not updated simultaneously.
  - **Write Back:** The processor updates only its cache, and the main memory is updated later. This approach can lead to inconsistencies if the data in the cache is not written back to the main memory before another processor reads it.

#### Example Scenario:

- Consider a multiprocessor system with two processors, P1 and P2, both sharing the main memory. Let X be a shared data element.
  - **Initial State:** Both processors have consistent copies of X.
  - **Write Through Policy:** If P1 writes a new value X' to its cache, the main memory is immediately updated, but P2's cache may still hold the old value X.
  - **Write Back Policy:** If P1 writes X' to its cache, the main memory and P2's cache remain with the old value X until P1's cache line is written back to the main memory.

### Coherence Problems in Process Migration

#### Process Migration:

- When a process containing a shared variable migrates from one processor to another, the coherence of the variable's data can be affected.

- **Write Back Cache:** If a process migrates from processor 1 to processor 2, the data modified by the process may reside in the cache of processor 1 and not be updated in the main memory. Thus, processor 2 may read stale data from the main memory.
- **Write Through Cache:** Even with write-through caches, inconsistencies can occur if the cache is not properly invalidated or updated during the migration process.

### Example Scenario:

- A process with a shared variable X migrates from processor 1 to processor 2.
  - **Write Back Cache:** Processor 1 has updated X to X' in its cache. If the process migrates to processor 2, it may read the old value X from the main memory since X' has not been written back yet.
  - **Write Through Cache:** If the process migrates, the caches need to be carefully managed to ensure consistency, as direct updates to the main memory may still leave other caches with stale data.

## 7) What is Cache Coherence Problem? What are the different causes of cache inconsistencies?

**Cache Coherence Problem** refers to the issue in multiprocessor systems where multiple processors cache copies of the same memory location. When one processor updates its cache, the other processors may still have stale data, leading to inconsistencies across the system.

### Causes of Cache Inconsistencies

#### 1. Simultaneous Writes:

- When two or more processors attempt to write different values to the same memory location simultaneously, it leads to inconsistencies as each processor's cache may reflect a different value.

#### 2. Write-Back Caches:

- In a write-back cache system, the processor updates its cache and postpones writing the updated value to the main memory. Other processors may read stale data from the main memory if they do not have the latest cached value.

#### 3. Write-Through Caches:

- Even with write-through caches, where updates are immediately written to the main memory, inconsistencies can occur if other processors' caches are not updated simultaneously. The updated value in the main memory might not be propagated to all caches promptly.

#### 4. Cache Invalidation Delays:

- When a processor updates a cached value, it should ideally invalidate or update the corresponding entries in other caches. Delays or failures in this invalidation process can result in other processors using outdated data.

#### 5. **Non-Synchronized Access:**

- If multiple processors access and modify shared data without proper synchronization mechanisms (like locks or semaphores), it can lead to race conditions, where the final value of the shared data depends on the sequence of operations, causing inconsistencies.

#### 6. **Memory Reordering:**

- Modern processors often reorder memory operations to optimize performance. This can lead to situations where the order of reads and writes seen by different processors is inconsistent, resulting in stale or inconsistent data being accessed.

### 8) **Explain Context Switching Policies**

Context switching is essential for multitasking in operating systems, enabling a single processor to manage multiple processes by switching between them. Different multithreaded architectures employ various context-switching policies, which are crucial for maximizing processor efficiency and minimizing idle time. Here are four primary context-switching policies:

#### 1. **Switch on Cache Miss:**

- **Policy:** The context is preempted when it causes a cache miss.
- **Details:** The processor switches contexts only when it is certain that the current one will be delayed for a significant number of cycles due to a cache miss. The average interval between misses ( $R$ ) and the time required to satisfy the miss ( $L$ ) are key factors.

#### 2. **Switch on Every Load:**

- **Policy:** The context switches on every load, regardless of whether it causes a miss.
- **Details:** This policy allows switching on every load, independent of a cache miss. Here, the average interval between loads ( $R$ ) is considered. Blocking occurs only if the load results in a cache miss.

#### 3. **Switch on Every Instruction:**

- **Policy:** The context switches on every instruction, regardless of its type.

- **Details:** This policy makes successive instructions independent, which benefits pipelined execution. It ensures that the processor frequently switches contexts, aiming to utilize pipeline stages efficiently.

#### 4. Switch on Block of Instructions:

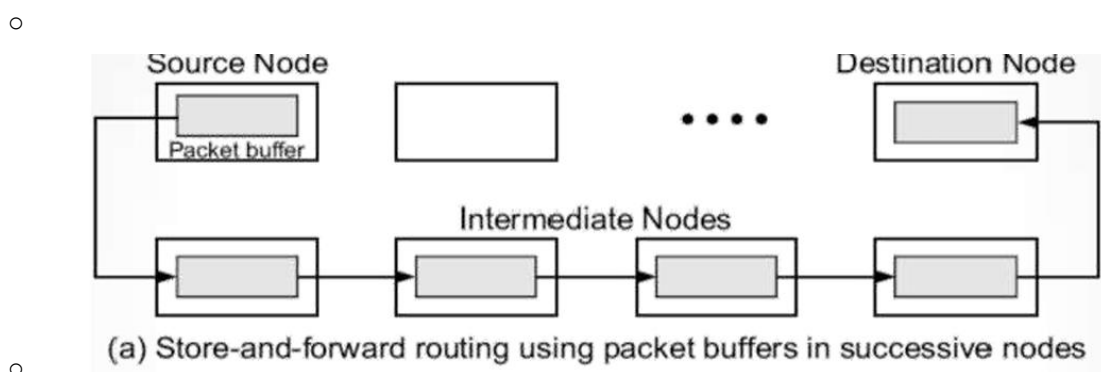
- **Policy:** Blocks of instructions from different threads are interleaved.
- **Details:** Interleaving blocks of instructions improves the cache-hit ratio due to locality and benefits single-context performance by making better use of cache memory n

### 9) Explain Store and forward routing wormhole routing related to message routing

#### Store and Forward Routing:

- **Mechanism:**

- In store and forward routing, packets are the basic unit of information flow.
- Each node in the network uses a packet buffer.
- A packet is transmitted from a source node to a destination node through a sequence of intermediate nodes.
- When a packet reaches an intermediate node, it is stored in the buffer first.
- The packet is forwarded to the next node only if the desired output channel and a packet buffer in the receiving node are both available.



- **Advantages:**

- Reliable as it ensures that packets are not lost during transmission.
- Simplifies the routing process since each node handles entire packets.

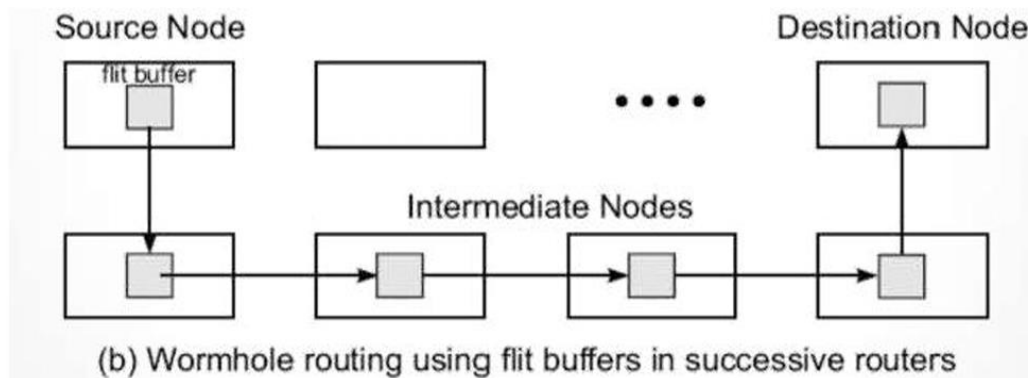
- **Disadvantages:**

- High latency due to the need to store and forward packets at each intermediate node.
- Requires large buffer sizes to store packets at each node.

#### Wormhole Routing:

- **Mechanism:**

- Packets are subdivided into smaller units called flits (flow control digits).
- Flit buffers are used in the hardware routers attached to nodes.
- Transmission from the source node to the destination node is done through a sequence of routers.
- Only the header flit knows the destination, and all data flits must follow the header flit.
- Flits from different packets cannot be mixed; otherwise, they may be routed to the wrong destination.
- Pipelined transmission: flits move asynchronously using a handshaking protocol.



- **Advantages:**

- Efficient as it reduces the need for large buffers.
- Faster transmission due to reduced latency.
- Suitable for high-performance computing environments.

- **Disadvantages:**

- Complex hardware requirements for managing flit buffers.
- Potential for deadlock if virtual channels are not managed properly.

### Comparison:

- **Latency:**

- Store and Forward Routing has a latency directly proportional to the distance between the source and the destination.
- Wormhole Routing has a latency almost independent of the distance, primarily dependent on the packet length and the flit length.

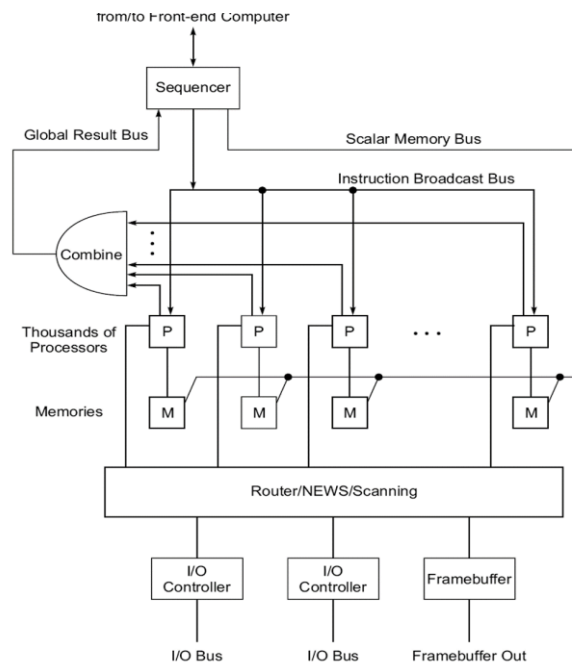
- **Buffer Requirements:**

- Store and Forward Routing requires larger buffers at each node to store entire packets.
- Wormhole Routing requires smaller flit buffers but needs sophisticated management to handle flit movement.

- **Performance:**

- Wormhole Routing generally offers better performance due to reduced latency and more efficient use of bandwidth .

## 10) Explain with diagram CM-2 Architecture



### Program Execution Paradigm

- **Execution Start:** All programs begin execution on a front-end.
- **Microinstructions:** The front-end issues microinstructions to the back-end processing array for data-parallel operations.
- **Sequencer:** These microinstructions are broken down by a sequencer and broadcast to all data processors in the array (Module-4 ACA).

### Data Exchange

Data sets and results can be exchanged between the front-end and the processing array in three ways:

1. **Broadcasting:** Carried out through a broadcast bus to all data processors at once.
2. **Global Combining:** Allows the front-end to obtain the sum, largest value, logical OR, etc., from values from each processor.
3. **Scalar Memory Bus:** Allows the front-end to read or write one 32-bit value at a time from or to the memories attached to the data processors (Module-4 ACA).

### Processing Array

- **Configuration:** The processing array contains from 4K to 64K bit-slice data processors (PEs).
- **Control:** All PEs are controlled by a sequencer (Module-4 ACA).

### Processing Nodes

- **Components:** Each data processing node contains 32 bit-slice data processors, an optional floating-point accelerator, and interfaces for inter-processor communication (Module-4 ACA).

### Hypercube Routers



- **Structure:** The router nodes on all processor chips are wired together to form a Boolean n-cube.
- **Configuration:** A full configuration of CM-2 has 4096 router nodes on processor chips interconnected as a 12-dimensional hypercube (Module-4 ACA).

### Major Applications

The CM-2 has been applied in various MPP and grand challenge applications, including:

- Document retrieval using relevance feedback
- Memory-based reasoning (e.g., medical diagnostic system called QUACK)
- Bulk processing of natural languages
- SPICE-like VLSI circuit analysis and layout
- Computational fluid dynamics

## 11) Explain Crossbar network with diagram?

### Crossbar Switch:

#### 1. Definition:

- A crossbar switch is a type of single-stage network.
- It is designed to connect every input to every output through a crosspoint switch.

#### 2. Structure:

- The crossbar switch uses a grid (or mesh) of switches where each switch connects a processor to a memory module.
- Each column in the  $n * m$  crossbar mesh contains  $n$  crosspoint switches.

#### 3. Functionality:

- Only one switch in each column can be connected at a time.
- The design must handle potential contention for each memory module.
- It avoids competition for bandwidth by using  $O(N^2)$  switches to connect  $N$  inputs to  $N$  outputs.

#### 4. Scalability:

- Crossbar switches are highly non-scalable due to their complexity and the number of switches required.
- They are typically used to connect a small number of workstations, generally 20 or fewer.

#### 5. Connection Process:

- Each processor provides a request line, a read/write line, a set of address lines, and a set of data lines to the crosspoint switch.
- The crosspoint switch responds with an acknowledgment once the access is completed.

